

## Assignment2 (Score: 10.0 / 12.0)

1. Test cell (Score: 2.0 / 2.0)
2. Test cell (Score: 2.0 / 2.0)
3. Test cell (Score: 1.0 / 1.0)
4. Test cell (Score: 2.0 / 2.0)
5. Test cell (Score: 2.0 / 2.0)
6. Test cell (Score: 0.0 / 2.0)
7. Test cell (Score: 1.0 / 1.0)

You are currently looking at **version 0.1** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.

# Assignment 2¶

In this assignment you'll explore the relationship between model complexity and generalization performance, by adjusting key parameters of various supervised learning models. Part 1 of this assignment will look at regression and Part 2 will look at classification.

## Part 1 - Regression¶

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

np.random.seed(0)
n = 15
x = np.linspace(0,10,n) + np.random.randn(n)/5
y = np.sin(x)+x/6 + np.random.randn(n)/10

X_train, X_test, y_train, y_test = train_test_split(x, y, random_state=0)

def intro():
    %matplotlib notebook

    plt.figure()
    plt.scatter(X_train, y_train, label='training data')
    plt.scatter(X_test, y_test, label='test data')
    plt.legend(loc=4);

intro()
```

## Question 1¶

Write a function that fits a polynomial LinearRegression model on the *training data* `X_train` for degrees 1, 3, 6, and 9. Use `PolynomialFeatures` in `sklearn.preprocessing` to create the polynomial features and then fit a linear regression model. For each model, find 100 predicted values over the interval  $x = 0$  to  $10$  (e.g. `np.linspace(0,10,100)`) and store this in a numpy array. The first row of this array should correspond to the output from the model trained on degree 1, the second row degree 3, the third row degree 6, and the fourth row degree 9.



The figure above shows the fitted models plotted on top of the original data (using `plot_one()`).

\*This function should return a numpy array with shape `(4, 100)`\*

In [2]:

```

def answer_one():
    from sklearn.linear_model import LinearRegression
    from sklearn.preprocessing import PolynomialFeatures

    # first step starts with creating an array with values between (1,10) using linspace w
    input_array = np.linspace(0, 10, 100)

    # We will create a new training data (X_train1) , with required shape (-1,1) which must
    # and also because, we will have the need of X_train for the next questions
    X_train1 = X_train.reshape(-1, 1)

    # I just created a 'predicted array' with a shape (4,100)
    # here, it is created with dtype='f8' which is float type because, if it is created in
    # we add float values to it, it will append them as int(that float number)
    # thats why, it is initially created as a float type array
    predicted_array = np.arange(400, dtype="f8").reshape(4, 100)

    # list of degrees that can be used for iteration
    degree = [1, 3, 6, 9]

    for x in range(4):
        # Now for each loop, we will create a new training data 'Xtrain_2' using 'X_train1'
        Xtrain_2 = PolynomialFeatures(degree=degree[x]).fit_transform(X_train1)

        # Now, we create a regressor that trains over 'Xtrain_2' and 'y_train'
        linreg = LinearRegression().fit(Xtrain_2, y_train)

        # now, we will predict the result of each regression using '.predict' method of the
        # be the polynomial featured data set for 'input_array'
        # here each row will be assigned with the predicted values of each degree
        predicted_array[x] = linreg.predict(
            PolynomialFeatures(degree=degree[x]).fit_transform(
                input_array.reshape(-1, 1)
            )
        )

    return predicted_array

answer_one()

```

Out[2]:

```

array([[ 2.53040195e-01,  2.69201547e-01,  2.85362899e-01,
         3.01524251e-01,  3.17685603e-01,  3.33846955e-01,
         3.50008306e-01,  3.66169658e-01,  3.82331010e-01,
         3.98492362e-01,  4.14653714e-01,  4.30815066e-01,
         4.46976417e-01,  4.63137769e-01,  4.79299121e-01,
         4.95460473e-01,  5.11621825e-01,  5.27783177e-01,
         5.43944529e-01,  5.60105880e-01,  5.76267232e-01,
         5.92428584e-01,  6.08589936e-01,  6.24751288e-01,
         6.40912640e-01,  6.57073992e-01,  6.73235343e-01,
         6.89396695e-01,  7.05558047e-01,  7.21719399e-01,
         7.37880751e-01,  7.54042103e-01,  7.70203454e-01,
         7.86364806e-01,  8.02526158e-01,  8.18687510e-01,

```

8.34848862e-01,	8.51010214e-01,	8.67171566e-01,
8.83332917e-01,	8.99494269e-01,	9.15655621e-01,
9.31816973e-01,	9.47978325e-01,	9.64139677e-01,
9.80301028e-01,	9.96462380e-01,	1.01262373e+00,
1.02878508e+00,	1.04494644e+00,	1.06110779e+00,
1.07726914e+00,	1.09343049e+00,	1.10959184e+00,
1.12575320e+00,	1.14191455e+00,	1.15807590e+00,
1.17423725e+00,	1.19039860e+00,	1.20655995e+00,
1.22272131e+00,	1.23888266e+00,	1.25504401e+00,
1.27120536e+00,	1.28736671e+00,	1.30352807e+00,
1.31968942e+00,	1.33585077e+00,	1.35201212e+00,
1.36817347e+00,	1.38433482e+00,	1.40049618e+00,
1.41665753e+00,	1.43281888e+00,	1.44898023e+00,
1.46514158e+00,	1.48130294e+00,	1.49746429e+00,
1.51362564e+00,	1.52978699e+00,	1.54594834e+00,
1.56210969e+00,	1.57827105e+00,	1.59443240e+00,
1.61059375e+00,	1.62675510e+00,	1.64291645e+00,
1.65907781e+00,	1.67523916e+00,	1.69140051e+00,
1.70756186e+00,	1.72372321e+00,	1.73988457e+00,
1.75604592e+00,	1.77220727e+00,	1.78836862e+00,
1.80452997e+00,	1.82069132e+00,	1.83685268e+00,
1.85301403e+00],		
[1.22989539e+00,	1.15143628e+00,	1.07722393e+00,
1.00717881e+00,	9.41221419e-01,	8.79272234e-01,
8.21251741e-01,	7.67080426e-01,	7.16678772e-01,
6.69967266e-01,	6.26866391e-01,	5.87296632e-01,
5.51178474e-01,	5.18432402e-01,	4.88978901e-01,
4.62738455e-01,	4.39631549e-01,	4.19578668e-01,
4.02500297e-01,	3.88316920e-01,	3.76949022e-01,
3.68317088e-01,	3.62341603e-01,	3.58943051e-01,
3.58041918e-01,	3.59558687e-01,	3.63413845e-01,
3.69527874e-01,	3.77821261e-01,	3.88214491e-01,
4.00628046e-01,	4.14982414e-01,	4.31198078e-01,
4.49195522e-01,	4.68895233e-01,	4.90217694e-01,
5.13083391e-01,	5.37412808e-01,	5.63126429e-01,
5.90144741e-01,	6.18388226e-01,	6.47777371e-01,
6.78232660e-01,	7.09674578e-01,	7.42023609e-01,
7.75200238e-01,	8.09124950e-01,	8.43718230e-01,
8.78900563e-01,	9.14592432e-01,	9.50714324e-01,
9.87186723e-01,	1.02393011e+00,	1.06086498e+00,
1.09791181e+00,	1.13499108e+00,	1.17202328e+00,
1.20892890e+00,	1.24562842e+00,	1.28204233e+00,
1.31809110e+00,	1.35369523e+00,	1.38877520e+00,
1.42325149e+00,	1.45704459e+00,	1.49007498e+00,
1.52226316e+00,	1.55352959e+00,	1.58379478e+00,
1.61297919e+00,	1.64100332e+00,	1.66778766e+00,
1.69325268e+00,	1.71731887e+00,	1.73990672e+00,
1.76093671e+00,	1.78032933e+00,	1.79800506e+00,
1.81388438e+00,	1.82788778e+00,	1.83993575e+00,
1.84994877e+00,	1.85784732e+00,	1.86355189e+00,
1.86698296e+00,	1.86806103e+00,	1.86670656e+00,
1.86284006e+00,	1.85638200e+00,	1.84725286e+00,
1.83537314e+00,	1.82066332e+00,	1.80304388e+00,
1.78243530e+00,	1.75875808e+00,	1.73193269e+00,
1.70187963e+00,	1.66851936e+00,	1.63177240e+00,
1.59155920e+00],		
[-1.99554310e-01,	-3.95192721e-03,	1.79851753e-01,
3.51005136e-01,	5.08831706e-01,	6.52819233e-01,
7.82609240e-01,	8.97986721e-01,	9.98870117e-01,

1.08530155e+00,	1.15743729e+00,	1.21553852e+00,
1.25996233e+00,	1.29115292e+00,	1.30963316e+00,
1.31599632e+00,	1.31089811e+00,	1.29504889e+00,
1.26920626e+00,	1.23416782e+00,	1.19076415e+00,
1.13985218e+00,	1.08230867e+00,	1.01902405e+00,
9.50896441e-01,	8.78825970e-01,	8.03709344e-01,
7.26434655e-01,	6.47876457e-01,	5.68891088e-01,
4.90312256e-01,	4.12946874e-01,	3.37571147e-01,
2.64926923e-01,	1.95718291e-01,	1.30608438e-01,
7.02167560e-02,	1.51162118e-02,	-3.41690365e-02,
-7.71657635e-02,	-1.13453547e-01,	-1.42666382e-01,
-1.64494044e-01,	-1.78683194e-01,	-1.85038228e-01,
-1.83421873e-01,	-1.73755533e-01,	-1.56019368e-01,
-1.30252132e-01,	-9.65507462e-02,	-5.50696231e-02,
-6.01973195e-03,	5.03325883e-02,	1.13667071e-01,
1.83611221e-01,	2.59742264e-01,	3.41589357e-01,
4.28636046e-01,	5.20322987e-01,	6.16050916e-01,
7.15183874e-01,	8.17052690e-01,	9.20958717e-01,
1.02617782e+00,	1.13196463e+00,	1.23755703e+00,
1.34218093e+00,	1.44505526e+00,	1.54539723e+00,
1.64242789e+00,	1.73537785e+00,	1.82349336e+00,
1.90604254e+00,	1.98232198e+00,	2.05166348e+00,
2.11344114e+00,	2.16707864e+00,	2.21205680e+00,
2.24792141e+00,	2.27429129e+00,	2.29086658e+00,
2.29743739e+00,	2.29389257e+00,	2.28022881e+00,
2.25656001e+00,	2.22312684e+00,	2.18030664e+00,
2.12862347e+00,	2.06875850e+00,	2.00156065e+00,
1.92805743e+00,	1.84946605e+00,	1.76720485e+00,
1.68290491e+00,	1.59842194e+00,	1.51584842e+00,
1.43752602e+00,	1.36605824e+00,	1.30432333e+00,
1.25548743e+00],		
[ 6.79501877e+00,	4.14319714e+00,	2.23123195e+00,
9.10495039e-01,	5.49803327e-02,	-4.41344174e-01,
-6.66950030e-01,	-6.94942449e-01,	-5.85049217e-01,
-3.85418102e-01,	-1.34235851e-01,	1.38818670e-01,
4.11275215e-01,	6.66715368e-01,	8.93747315e-01,
1.08510182e+00,	1.23683955e+00,	1.34766042e+00,
1.41830603e+00,	1.45104693e+00,	1.44924662e+00,
1.41699501e+00,	1.35880409e+00,	1.27935949e+00,
1.18332142e+00,	1.07516952e+00,	9.59085935e-01,
8.38871928e-01,	7.17893068e-01,	5.99048939e-01,
4.84763319e-01,	3.76991254e-01,	2.77239711e-01,
1.86598853e-01,	1.05781227e-01,	3.51664578e-02,
-2.51506700e-02,	-7.53106421e-02,	-1.15639771e-01,
-1.46602278e-01,	-1.68755085e-01,	-1.82706256e-01,
-1.89077879e-01,	-1.88473948e-01,	-1.81453660e-01,
-1.68510358e-01,	-1.50056232e-01,	-1.26412707e-01,
-9.78063691e-02,	-6.43701362e-02,	-2.61492812e-02,
1.68881553e-02,	6.48371253e-02,	1.17838121e-01,
1.76057180e-01,	2.39664066e-01,	3.08809354e-01,
3.83601193e-01,	4.64082501e-01,	5.50209339e-01,
6.41831222e-01,	7.38674048e-01,	8.40326319e-01,
9.46229255e-01,	1.05567134e+00,	1.16778775e+00,
1.28156502e+00,	1.39585128e+00,	1.50937206e+00,
1.62075184e+00,	1.72854110e+00,	1.83124870e+00,
1.92737900e+00,	2.01547327e+00,	2.09415450e+00,
2.16217452e+00,	2.21846241e+00,	2.26217255e+00,
2.29273075e+00,	2.30987650e+00,	2.31369910e+00,
2.30466527e+00,	2.28363544e+00,	2.25186569e+00,

```
2.21099194e+00, 2.16299281e+00, 2.11012698e+00,  
2.05484079e+00, 1.99964137e+00, 1.94693015e+00,  
1.89879129e+00, 1.85672916e+00, 1.82134864e+00,  
1.79197149e+00, 1.76618168e+00, 1.73929211e+00,  
1.70372475e+00, 1.64829557e+00, 1.55739548e+00,  
1.41005768e+00]])
```

In [3]:

Grade cell: cell-4b3a4b2c2971710c

Score: 2.0 / 2.0

```
# feel free to use the function plot_one() to replicate the figure  
# from the prompt once you have completed question one  
def plot_one(degree_predictions):  
    plt.figure(figsize=(10,5))  
    plt.plot(X_train, y_train, 'o', label='training data', markersize=10)  
    plt.plot(X_test, y_test, 'o', label='test data', markersize=10)  
    for i,degree in enumerate([1,3,6,9]):  
        plt.plot(np.linspace(0,10,100), degree_predictions[i], alpha=0.8, lw=2, label='degr  
    plt.ylim(-1,2.5)  
    plt.legend(loc=4)  
  
plot_one(answer_one())
```

## Question 21

Write a function that fits a polynomial LinearRegression model on the training data `X_train` for degrees 0 through 9. For each model compute the  $R^2$  (coefficient of determination) regression score on the training data as well as the test data and return both of these arrays in a tuple.

This function should return a tuple of numpy arrays (`r2_train`, `r2_test`). Both arrays should have shape `(10,)`

In [4]:

Student's answer

```

def answer_two():
    from sklearn.linear_model import LinearRegression
    from sklearn.metrics import r2_score
    from sklearn.preprocessing import PolynomialFeatures

    # create empty lists for r2_train and r2_test
    r2_train = []
    r2_test = []

    # then, create new data sets based on X_train and X_test with the required shape of onl
    X_train1 = X_train.reshape(-1, 1)
    X_test1 = X_test.reshape(-1, 1)

    for x in range(10):
        # Now, create polynomial featured datasets from X_train1, X_test1 which is used to
        X_train2 = PolynomialFeatures(degree=x).fit_transform(X_train1)
        X_test2 = PolynomialFeatures(degree=x).fit_transform(X_test1)

        # Now, train the regressor with the above obtained polynomial featured training dat
        linreg = LinearRegression().fit(X_train2, y_train)

        y_predicted_train = linreg.predict(X_train2)
        y_predicted_test = linreg.predict(X_test2)

        # Now, now append the scores of training and testing data sets into the respective
        r2_train.append(r2_score(y_train, y_predicted_train))
        r2_test.append(r2_score(y_test, y_predicted_test))

    # Now, convert the lists into arrays of required shape(10,1)
    # Note - in this question, it is asked to make the shape of array as (10,1) but auto gr
    # reshape it as (1,10)

    r2_train = np.array(r2_train).reshape(10, 1)
    r2_test = np.array(r2_test).reshape(10, 1)

    # Your code here
    ans = tuple([r2_train, r2_test])

    return ans

answer_two()

```

Out[4]:

```
(array([[0.         ],
       [0.42924578],
       [0.4510998 ],
       [0.58719954],
       [0.91941945],
       [0.97578641],
       [0.99018233],
       [0.99352509],
       [0.99637545],
       [0.99803706]]),
 array([[-0.47808642],
       [-0.45237104],
       [-0.06856984],
       [ 0.00533105],
       [ 0.73004943],
       [ 0.87708301],
       [ 0.9214094 ],
       [ 0.92021504],
       [ 0.63247942],
       [-0.64525285]]))
```

In [5]:

Grade cell: cell-84a9de5d00b87d2f

Score: 2.0 / 2.0

### Question 3

Based on the  $R^2$  scores from question 2 (degree levels 0 through 9), what degree level corresponds to a model that is underfitting? What degree level corresponds to a model that is overfitting? What choice of degree level would provide a model with good generalization performance on this dataset?

(Hint: Try plotting the  $R^2$  scores from question 2 to visualize the relationship)

*This function should return a tuple with the degree values in this order: (Underfitting, Overfitting, Good\_Generalization)*



In [6]:

Student's answer

```
def answer_three():  
    # Your code here  
  
    (r2_train, r2_test) = answer_two()  
    # print(r2_test)  
    # print(r2_train)  
    # import matplotlib.pyplot as plt  
    # %matplotlib notebook  
    # plt.figure()  
    # plt.plot(r2_train, label='data')  
    # plt.plot(r2_test, label='data')  
    order = (2,9,7)  
    return (order) # Return  
answer_three()
```

Out[6]:

(2, 9, 7)

In [7]:

Grade cell: cell-877e9f32963e0d5a

Score: 1.0 / 1.0

## Question 4

Training models on high degree polynomial features can result in overfitting. Train two models: a non-regularized LinearRegression model and a Lasso Regression model (with parameters  $\alpha=0.01$ ,  $\text{max\_iter}=10000$ ,  $\text{tol}=0.1$ ) on polynomial features of degree 12. Return the  $R^2$  score for LinearRegression and Lasso model's test sets.

This function should return a tuple (LinearRegression\_R2\_test\_score, Lasso\_R2\_test\_score)

In [8]:

Student's answer

```
def answer_four():
    from sklearn.preprocessing import PolynomialFeatures
    from sklearn.linear_model import Lasso, LinearRegression
    from sklearn.metrics import r2_score
    from sklearn.preprocessing import MinMaxScaler
    scaler = MinMaxScaler()

    poly = PolynomialFeatures(degree=12)
    X_poly = poly.fit_transform(X_train.reshape(11,1))
    X_test_poly = poly.fit_transform(X_test.reshape(4,1))

    # Create regressor
    linreg = LinearRegression().fit(X_poly, y_train)

    X_train_scaled = scaler.fit_transform(X_poly)
    X_test_scaled = scaler.transform(X_test_poly)

    linlasso = Lasso(alpha=0.01, max_iter = 10000, tol = 0.1).fit(X_poly, y_train)

    return r2_score(y_test, linreg.predict(X_test_poly)) , r2_score(y_test, linlasso.predict(X_test_scaled))

answer_four()
```

Out[8]:

```
(-4.3119675863308435, 0.6051396919570036)
```

In [9]:

Grade cell: cell-f1ccfec3713e840c

Score: 2.0 / 2.0

## Part 2 - Classification¶

For this section of the assignment we will be working with the UCI Mushroom Data Set (<http://archive.ics.uci.edu/ml/datasets/Mushroom?ref=datanews.io>) stored in `mushrooms.csv`. The data will be used to train a model to predict whether or not a mushroom is poisonous. The following attributes are provided:

*Attribute Information:*

1. cap-shape: bell=b, conical=c, convex=x, flat=f, knobbed=k, sunken=s
2. cap-surface: fibrous=f, grooves=g, scaly=y, smooth=s
3. cap-color: brown=n, buff=b, cinnamon=c, gray=g, green=r, pink=p, purple=u, red=e, white=w, yellow=y
4. bruises?: bruises=t, no=f
5. odor: almond=a, anise=l, creosote=c, fishy=y, foul=f, musty=m, none=n, pungent=p, spicy=s
6. gill-attachment: attached=a, descending=d, free=f, notched=n
7. gill-spacing: close=c, crowded=w, distant=d
8. gill-size: broad=b, narrow=n
9. gill-color: black=k, brown=n, buff=b, chocolate=h, gray=g, green=r, orange=o, pink=p, purple=u, red=e, white=w, yellow=y
10. stalk-shape: enlarging=e, tapering=t
11. stalk-root: bulbous=b, club=c, cup=u, equal=e, rhizomorphs=z, rooted=r, missing=?
12. stalk-surface-above-ring: fibrous=f, scaly=y, silky=k, smooth=s
13. stalk-surface-below-ring: fibrous=f, scaly=y, silky=k, smooth=s
14. stalk-color-above-ring: brown=n, buff=b, cinnamon=c, gray=g, orange=o, pink=p, red=e, white=w, yellow=y
15. stalk-color-below-ring: brown=n, buff=b, cinnamon=c, gray=g, orange=o, pink=p, red=e, white=w, yellow=y
16. veil-type: partial=p, universal=u
17. veil-color: brown=n, orange=o, white=w, yellow=y
18. ring-number: none=n, one=o, two=t
19. ring-type: cobwebby=c, evanescent=e, flaring=f, large=l, none=n, pendant=p, sheathing=s, zone=z
20. spore-print-color: black=k, brown=n, buff=b, chocolate=h, green=r, orange=o, purple=u, white=w, yellow=y
21. population: abundant=a, clustered=c, numerous=n, scattered=s, several=v, solitary=y
22. habitat: grasses=g, leaves=l, meadows=m, paths=p, urban=u, waste=w, woods=d

The data in the mushrooms dataset is currently encoded with strings. These values will need to be encoded to numeric work with sklearn. We'll use `pd.get_dummies` to convert the categorical variables into indicator variables.

In [10]:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split

mush_df = pd.read_csv('assets/mushrooms.csv')
mush_df2 = pd.get_dummies(mush_df)

X_mush = mush_df2.iloc[:,2:]
y_mush = mush_df2.iloc[:,1]

X_train2, X_test2, y_train2, y_test2 = train_test_split(X_mush, y_mush, random_state=0)
```

In [11]:

```
X_subset = X_test2  
y_subset = y_test2
```

## Question 5

Using `X_train` and `y_train` from the preceeding cell, train a `DecisionTreeClassifier` with default parameters and `random_state=0`. What are the 5 most important features found by the decision tree?

*This function should return a list of length 5 of the feature names in descending order of importance.*

In [12]:

Student's answer

```
def answer_five():  
    from sklearn.tree import DecisionTreeClassifier  
  
    # First, creat a classifer using the given data sets X_train and y_train2  
    classifier = DecisionTreeClassifier(random_state=0).fit(X_train2, y_train2)  
  
    # then, using the feature_importance attribute, create a list of those feature importan  
    features_list = list(classifier.feature_importances_)  
    features_list.sort(reverse=True)  
    top_5_features_num_list = features_list[:5]  
  
    # then, find those index positions, where the top 5 featureres are in the list  
    index_pos_list = [  
        list(classifier.feature_importances_).index(x) for x in top_5_features_num_list  
    ]  
  
    # then, find the name of the features from the .columns of the dataset X_train2  
    top_5_features_list = [X_train2.columns[x] for x in index_pos_list]  
    # Your code here  
  
    return top_5_features_list  
  
answer_five()
```

Out[12]:

```
['odor_n', 'stalk-root_c', 'stalk-root_r', 'spore-print-color_r', 'odor_l']
```

In [13]:

Grade cell: cell-45bb9fef0723e714

Score: 2.0 / 2.0

## Question 6

For this question, use the `validation_curve` function in `sklearn.model_selection` to determine training and test scores for a Support Vector Classifier (SVC) with varying parameter values.

Create an SVC with default parameters (i.e. `kernel='rbf'`, `C=1`) and `random_state=0`. Recall that the kernel width of the RBF kernel is controlled using the `gamma` parameter. Explore the effect of `gamma` on classifier accuracy by using the `validation_curve` function to find the training and test scores for 6 values of `gamma` from `0.0001` to `10` (i.e. `np.logspace(-4,1,6)`).

For each level of `gamma`, `validation_curve` will use 3-fold cross validation (use `cv=3`, `n_jobs=2` as parameters for `validation_curve`), returning two 6x3 (6 levels of `gamma` x 3 fits per level) arrays of the scores for the training and test sets in each fold.

Find the mean score across the five models for each level of `gamma` for both arrays, creating two arrays of length 6, and return a tuple with the two arrays.

e.g.

if one of your array of scores is

```
array([[ 0.5,  0.4,  0.6],
       [ 0.7,  0.8,  0.7],
       [ 0.9,  0.8,  0.8],
       [ 0.8,  0.7,  0.8],
       [ 0.7,  0.6,  0.6],
       [ 0.4,  0.6,  0.5]])
```

it should then become

```
array([ 0.5,  0.73333333,  0.83333333,  0.76666667,  0.63333333,  0.5])
```

*This function should return a tuple of numpy arrays (`training_scores`, `test_scores`) where each array in the tuple has shape (6,).*

In [14]:

Student's answer

```
def answer_six():
    from sklearn.svm import SVC
    from sklearn.model_selection import validation_curve

    # Your code here
    this_C = 1.0
    clf = SVC(kernel = 'rbf', C=this_C).fit(X_train2, y_train2)

    param_range = np.logspace(-4,1,6)
    # print(C)

    train_scores, test_scores = validation_curve(clf, X_subset, y_subset, param_name='gamma'
                                                param_range=param_range, cv=3)
    training_scores_mean = np.mean(train_scores, axis = 1)
    test_scores_mean = np.mean(test_scores, axis = 1)

    return (training_scores_mean, test_scores_mean)

answer_six()
```

Out[14]:

```
(array([0.56646972, 0.93106844, 0.990645 , 1. , 1. , 1. ],
      [1. , 1. , 1. , 1. , 1. , 1. ]),
 array([0.56720827, 0.9300837 , 0.98966027, 1. , 1. , 0.99458395,
        0.52240276]))
```

In [15]:

Grade cell: cell-ae2a630e8862c3a3

Score: 0.0 / 2.0

You have failed this test due to an error. The traceback has been removed because it may contain sensitive information.

AssertionError: Q6: The training\_scores has incorrect value at index 0

## Question 7

Based on the scores from question 6, what gamma value corresponds to a model that is underfitting? What gamma value corresponds to a model that is overfitting? What choice of gamma would provide a model with good generalization performance on this dataset?

(Hint: Try plotting the scores from question 6 to visualize the relationship)

*This function should return a tuple with the degree values in this order: (Underfitting, Overfitting, Good\_Generalization)*

In [16]:

Student's answer

```
def answer_seven():  
    #As, there are limited number of elements, we can conclude things, just by looking over the  
  
    #For overfitting, i think gamma = 10^1 is the correct as, it has the best training score and  
  
    #For underfitting, i think gamma = gamma = 10^-4 is the correct one, as it has least training  
  
    #For good generalization, gamma = 10^-1 is the best, as it has a training score and testing  
    # Your code here  
    ans = tuple([0.0001,10,0.1])  
    return ans  
answer_seven()
```

Out[16]:

(0.0001, 10, 0.1)

In [17]:

Grade cell: cell-6372c6701a14c068

Score: 1.0 / 1.0

In [ ]:

In [ ]:

This assignment was graded by mooc\_adswpy:e5e20d3b91dd, v1.45.052423